

An Introduction to one-class classification with the oneClass Package

Benjamin Mack

June 26, 2014

Contents

1	One-class classification	2
2	Installation	3
3	Synthetic data: bananas	3
4	Model selection in the absence of PN-data.	4
5	Evaluation with PN-data	10
6	More one-class classifiers...	12
7	Parallel processing	12
8	Prediction of large raster data	13
9	Future work ...	13

1 One-class classification

The purpose of **one-class classifier** is identical to the to classify new data based on a classification model learned from labeled examples. For the training of the classifier labeled samples are only required for the positive class, i.e. the class of interest. Here is the difference to a binary classifier which has to be trained with labeled data of both the positive and the negative class. Thus, a one-class classifier is employed when the acquisition of representative labeled data for the negative class is expensive or not possible at all.

The convenience of not requiring negative training data comes at a price because OCC challenging due to the limited information contained in the training set. Unlabeled training data can, and usually should, be used to build more accurate predictive models. However, the process is still uncertain and the classification outcome has to be treated with caution.

A paragraph about the importance of model selection, the difficulty when only PU-data is available, about resampling and performance metrics

The package `oneClass` shall serve the requirements of **two potential users**, the *analyst* and the *developer*. These are extrem characters and in reality one will usually be located somewhere in between.

The *analyst* is faced with a particular one-class classification problem, i.e. a set of positive training samples and the unlabeled data to be classified. It is assumed that no complete and representative test set is available for the purpose of validation and testing. In such a situation a careful evaluation of the classification outcome based on the available (positive and unlabeled) data is required in order to select the most promising final model. In section4 the analysis strategy presented in [REF](#) is followed and the main functions are illustrated. Hopefully the package can be helpful to solve a given one-class classification problem more effective and more convenient in practice.

The *developer* is interested in the developement of new and/or optimization of existing methods. Therefore the package provides convenient functionalities. The package `oneClass` builds upon the powerful package `caret` [1] and tries to adapt the philosophy.

The package `caret` allows one to embed own **custom functionalities** in the rich infrastructure of the `caret` package, in fact the one-class classification methods of the `oneClass` are implemented as custom functions which are passed to the function `train()`. This infrastructure comprises methods for pre-processing data, calculating variable importance, model visualizations, and parallel processing.

The function `oneClass()` calls the `caret::train()` function with a one-class methods. It returns an object of class `oneClass` which inherits from the class `train`. Thus, many of the powerful analysis tools from the `train` package can also be used seamlessly when using an object of class `oneClass`.

One-class classifiers

The `oneClass` package is a user-oriented environment for analyzing one-class classification problems. It implements two commonly used classifiers, the biased SVM [2, 3], and Maxent [4, 5].

These classifiers are implemented as costum functions for the `train()` function from the `caret` package [\[REF\]](#). Thus the extensive functionalitie of `train()` can be used for model selection.

PU-performance metrics

...




Figure 1: The class membership (left), and the two features (middle and right) of the synthetic bananas data set.

2 Installation

Currently the package is still not on R CRAN but can be downloaded from GitHub (<https://github.com/benmack/oneClass>). It can be installed in R with the package devtools and the following commands (This may take a while...):

```
library(devtools)
install_github('benmack/oneClass')
```

3 Synthetic data: bananas

In the following the package is demonstrated by means of the synthetic banana data set which is included in the package.

The bananas data is stored as raster data, where `y` is a one-band raster with the class patches, i.e. what we want to find out when performing one-class classification with remotely sensed data. `x` are the features or predictors based on which the classification has to be build.

```
require(oneClass)
require(raster)
data(bananas)
plot(stack(bananas$y, bananas$x), nc=3)
```

In one-class classification a set of positive labeled samples is available for training the classifier. Furthermore, unlabeled samples are used, which are usually a random sample of the whole data. Such a training data set is also stored in the bananas data set. Additionally we generate a test data set consisting of 1.000 random samples of the image. As we mentioned before, we would not expect to have a test set in real-world one-class classification application. But here we also create a PN test data set (PN-data) to show the available (PN-) evaluation methods which are useful when investigating one-class classifiers of developing new methods.

```

seed <- 123456
tr.x <- bananas$tr[, -1]
tr.y <- puFactor(bananas$tr[, 1], positive=1)
set.seed(seed)
te.i <- sample(ncell(bananas$y), 1000)
te.x <- extract(bananas$x, te.i)
te.y <- extract(bananas$y, te.i)


```

70 With the two-dimensional synthetic data set we can view the data in the feature space:

```

par(mfrow=c(1,2), mgp=c(1.75, .75, 0), mar=c(3, 3, .1, .1))
plot(tr.x, pch=ifelse(tr.y==1, 16, 4) )
plot(te.x, pch=ifelse(te.y==1, 16, 4) )

```



71

72 A one-class classifier is supposed to learn from the PU-data (left) to optimally classify the PN data (right).

73 4 Model selection in the absence of PN-data.

74 `oneClass()` requires the training data as input. We also pass the whole unlabeled data because we want to
 75 analyze the predicted outcome of the whole data to better examine the performance of the model. Of course,
 76 if the data is very large, it is possible to use a subset here.

```

set.seed(seed)

```

```

index <- createFolds( tr.y, k=10, returnTrain=TRUE )
trControl <- trainControl(method='cv',
                           index=index,
                           summaryFunction = puSummary, # PU-performance metrics
                           classProbs=TRUE,             # important
                           savePredictions = TRUE,        # important
                           returnResamp = 'all')          # for resamples.train
oc <- oneClass( x = tr.x, y = tr.y, trControl = trControl )
pred <- predict(oc, bananas$x)

```


In a later section the internal processing steps of `oneClass()` are explained in more detail. Now we want to understand if the trained model is reasonable. The distributions of the predicted training data (`tr.x`) and unlabeled data (`bananas$x`) can help to make a first diagnosis. They are figures in the left plot below which we refer to as the 'diagnostic plot' (histogram: predicted unlabeled data, blue/grey boxplots: held-out predictions of the positive/unlabeled training samples). The right plot shows the model, i.e. the separating hyperplane of the biased SVM model (black line) and the distances (color coded). Please note that the color code in the right plot corresponds to the x-axis in the left plot.

Of course, in practice the right plot can not be visualized because the input data space is usually high dimensional. We show it here to facilitate the understanding of the interpretation.

```

plot(pred)
hist(oc, pred)
featurespace(oc, th=0)

```



From the diagnostic plot important conclusions can be derived about the training data, the class separability, and the suitability of a threshold for binarization, i.e. the conversion of the continuous classifier output in a binary classification.

In the envisaged strategy, the careful interpretation of the diagnostic plot is a critical element of the whole one-class classification processing chain. It provides the analyst with information based on which decisions are made.

The interpretation of the diagnostic plot reads as follows:


Regarding the **training data** we can assume that the amount of unlabeled training data is sufficient. This is because the positive and unlabeled training data overlaps (recognizable by the boxplots). This is very important. Imagine the complete unlabeled training data would be located at very low predictive values. In this case we could not be confident that the unlabeled training contains the relevant information required to build the model. Of course, if the data is well separable and/or the optimal decision boundary is very complex (non-linear) it is still possible to derive a good classifier. However, if this is not the case, the model is likely to be biased and/or to underfit the data.

Regarding the **separability** the following conclusions can be made:
The positive training data corresponds well with a distinctive cluster of data in the histogram at high predictive values. This cluster is separated by a low density area from the rest of the majority of the data. If we assume that the positive data set is representative and sufficiently large for the positive class we can conclude that the classification model is suitable and has a high discriminative power.


Regarding the **suitability of the threshold** we can assume that any threshold in the low density region leads to relatively high accuracies.

Nevertheless, it is always a good idea to examine the relationship between the estimated PU-performance and the tuning parameters. caret offers different plotting functions for this purpose¹. For example, here the puF and puAuc performance metrics are plotted in a heatmap.

```
trellis.par.set(caretTheme()) # nice colors from caret
plot(oc, metric = 'puF', plotType = "level")
plot(oc, metric = 'puAuc', plotType = "level")
```



¹<http://caret.r-forge.r-project.org/training.html>



These plots are informative to evaluate the tuning process and the final model. First, we can see if the parameter space in which we were looking for the optimal model has been reasonable. In the case here this is true because the models with the highest PU-performances are not located at the border of the parameter space. If this would be the case, a more suitable parameter combination could possibly lead to more powerful model. Also in the case where the optimal models are located at intermediate parameter values it is possible to find better models by repeating the model tuning over a finer grid of parameters in the space of optimal models. This is particularly recommendable if the estimated performance does not change smoothly with the parameters.

We can also examine the table which is printed when printing an object of class `train` or its inheriting child `oneClass`. Here we do not print the full table, as it is too large but an ordered and subsetting version of the complete table.

```
# oc # --> prints the whole large table with performance metrics of all models
sort(oc, printTable=TRUE, rows=1:10, by='puF', digits=2)

##      sigma      cNeg cMultiplier puAuc puF  Tpr puPpv puAucSD puFSD TprSD puPpvSD
## 30    0.1 4.0e+00         16 0.92 6.1 0.95 0.16 0.107 2.1 0.16 0.040
## 38    0.1 3.2e+01         64 0.92 5.4 0.90 0.16 0.091 2.1 0.21 0.054
## 37    0.1 3.2e+01         16 0.91 5.4 0.85 0.15 0.117 2.8 0.24 0.044
## 79    1.0 6.2e-02         16 0.91 5.3 0.90 0.16 0.082 2.1 0.21 0.041
## 80    1.0 6.2e-02         64 0.91 5.1 0.95 0.19 0.086 1.8 0.16 0.049
## 86    1.0 5.0e-01         16 0.90 5.1 0.85 0.15 0.099 2.0 0.24 0.043
## 93    1.0 4.0e+00         16 0.85 5.0 0.65 0.12 0.154 3.9 0.41 0.032
## 31    0.1 4.0e+00         64 0.93 4.9 0.95 0.20 0.090 2.0 0.16 0.060
## 44    0.1 2.6e+02         16 0.89 4.8 0.80 0.14 0.117 2.4 0.26 0.049
## 51    0.1 2.0e+03         16 0.85 4.4 0.65 0.13 0.144 2.9 0.41 0.047
```

```
sort(oc, printTable=TRUE, rows=1:10, by='puAuc', digits=2)
```

##	sigma	cNeg	cMultiplier	puAuc	puF	Tpr	puPpv	puAucSD	puFSD	TprSD	puPpvSD	
## 31	0.1	4.0000		64	0.93	4.9	0.95	0.204	0.090	2.0	0.16	0.060
## 30	0.1	4.0000		16	0.92	6.1	0.95	0.158	0.107	2.1	0.16	0.040
## 38	0.1	32.0000		64	0.92	5.4	0.90	0.160	0.091	2.1	0.21	0.054
## 71	1.0	0.0078		4	0.92	0.0	0.00	0.000	0.050	0.0	0.00	0.000
## 24	0.1	0.5000		64	0.91	3.4	0.95	0.287	0.105	1.4	0.16	0.067
## 36	0.1	32.0000		4	0.91	3.1	0.40	0.071	0.110	2.3	0.32	0.035
## 37	0.1	32.0000		16	0.91	5.4	0.85	0.148	0.117	2.8	0.24	0.044
## 79	1.0	0.0625		16	0.91	5.3	0.90	0.162	0.082	2.1	0.21	0.041
## 78	1.0	0.0625		4	0.91	0.0	0.00	0.000	0.065	0.0	0.00	0.000
## 80	1.0	0.0625		64	0.91	5.1	0.95	0.188	0.086	1.8	0.16	0.049

Thus, according to the puF-performance metric the model with the parameters sigma=0.1, cNeg=100, and cMultiplier=8 is the best. But other models are close behind the metric. Furthermore, there are other performance metrics, e.g. the puAuc which do not agree with the ranking of the puF.

It is difficult to say which metric is better. They both have different characteristics. The f1Pu is a measure which only evaluates at the default threshold 0. Thus, even models with high discriminative power can be placed at a low rank when the threshold 0 is unsuitable. The puAuc instead measures the discriminative over the whole range of possible threshold values. However, this can also introduce misleading measures, due to the fact that thresholds influence the measure which are far from being rational choices [REF].

Inspection of the resampling distributions of different models can give more evidence on which the model selection can be founded. Let us first find ten models which are ranked highest by the puF and/or the puAuc.

```

puF.ranking <- sort (oc, by='puF', print=FALSE)
puAuc.ranking <- sort (oc, by='puAuc', print=FALSE)

candidates <- as.numeric(unique(c(rownames(puF.ranking),
                                rownames(puAuc.ranking)) [
                                rep(1:10, each=2)+rep(c(0, 10), 10) ]))

candidates[1:10]

## [1] 30 23 38 29 37 58 79 87 80 88


```

For these models let us extract the performance metrics of the resamples and compare the resampling distributions between the models.

```

resamps <- resamples.train(oc, modParam=oc$results[candidates,1:3])
bwplot(resamps, metric='puF')

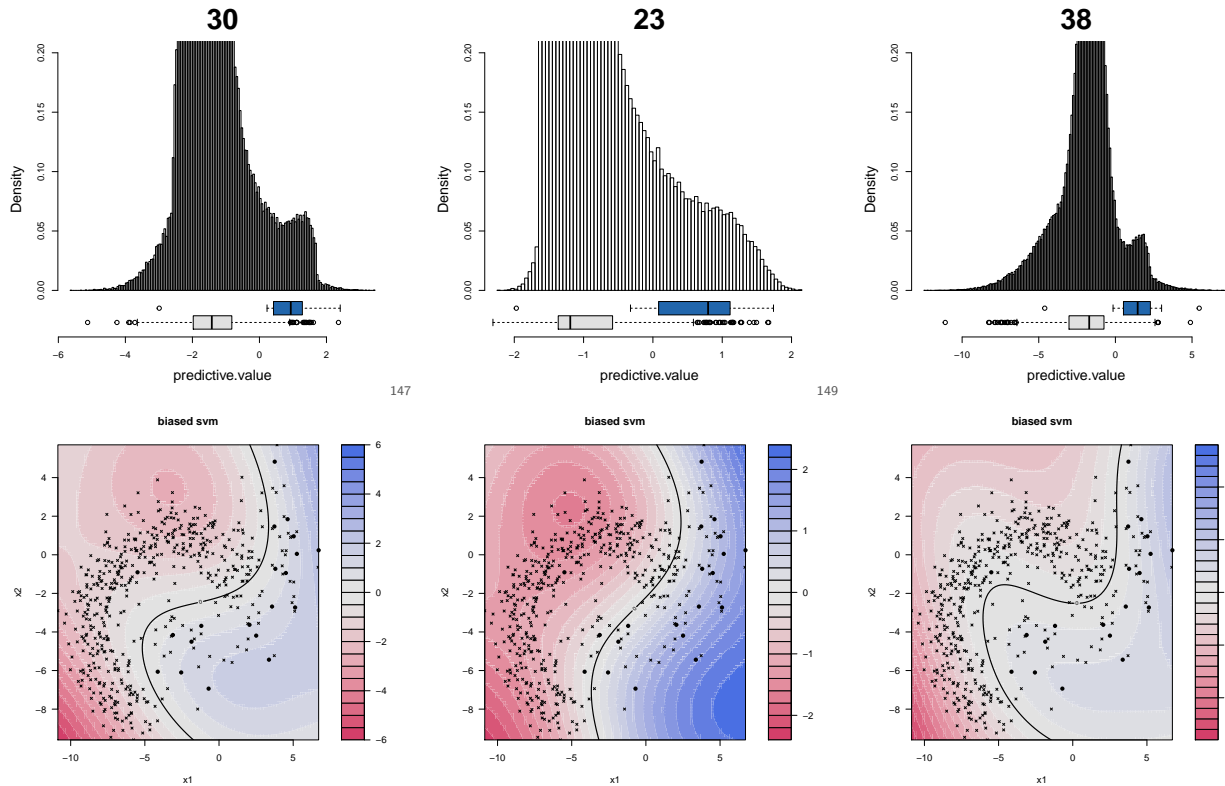
```

139

140 As we can see that the model ranked highest by the puF measure (second boxplot from the top) has com-
 141 peting models with very similar and resampling distributions. These models are not significantly different
 142 in terms of the resampling distributions so any one of them is a potential final model. It can be helpful to
 143 investigate the diagnostic histogram plot of the whole unlabeled data in order to gain more insight in the
 144 model characteristics.

```
oc.candidate1 <- update(oc, u=bananas$x, modRow=candidates[1])
oc.candidate2 <- update(oc, u=bananas$x, modRow=candidates[2])
oc.candidate3 <- update(oc, u=bananas$x, modRow=candidates[3])
```



Note that the model #35 (i.e. the model in the 35th row of the `oc$results` data frame) is ranked highest in the resamples-plot [Find out how the ranking works? Median? Can it be customized?](#). Second the diagnostic histogram plot shows a clear distinctive data cluster at high predictive values separated by a low density area from the large part of the data. From the diagnostic histogram plot we would intuitively identify the one of model #35 as the one which shows highest discriminative power, or in other words smallest overlapping between the positive and the negative class distributions. If we require a binary classification result we would intuitively threshold the predictive outcome at around 0, i.e. somewhere in the separating low density area.

Instead, in model 26 we should expect a higher overlap and it is less trivial to make a decision about a binarization threshold.

5 Evaluation with PN-data

In the last section we assumed that no complete and representative test data is available when the model has to be constructed. In other situations PN-data might be available and

However, in other many situations PN-data is available which makes a traditional accuracy assessment possible. There are different aspects we might want to evaluate. Given a particular model and threshold have been selected the confusion matrix and thereof derived accuracy measures for the binary classification result are of interest. We can calculate this with the function `evaluate`:

```
te.pred <- predict(oc, bananas$x[te.i])
```

```

ev <- dismo::evaluate(p=te.pred[te.y==1], a=te.pred[te.y!=1])
ev.th0 <- evaluate(ev, th=0)
ev.th0

## Positive/negative (+/-) test samples:
## 105 / 895
##
## Confusion Matrix at threshold 0.0084 :
##
##          + (Test) - (Test)  SUM UA[%]
## + (Pred)      75      17   92   82
## - (Pred)      30     878  908   97
## SUM          105     895 1000
## PA[%]         71      98
## ---
## OA[%]         95
## AUC[*100]     97
## K[*100]       74


```

If we are interested in the evolution of the accuracy over different thresholds we can also specify a vector of threshold or just leave the argument away in which case several thresholds are generated automatically over the range of the predictive values. From this we might also derive maximum achievable accuracy, i.e. the accuracy at the threshold which optimizes a given accuracy metric, e.g. the κ coefficient.

```

plot(ev)
evaluate(ev, th='max.kappa')

```



```

## Positive/negative (+/-) test samples:
## 105 / 895
##
## Confusion Matrix at threshold -0.1871 :
##
##          + (Test) - (Test)  SUM UA[%]
## + (Pred)      86      30  116   74
## - (Pred)      19     865  884   98
## SUM          105     895 1000
## PA[%]         82      97
## ---
## OA[%]         95
## AUC[*100]     97
## K[*100]       75

```

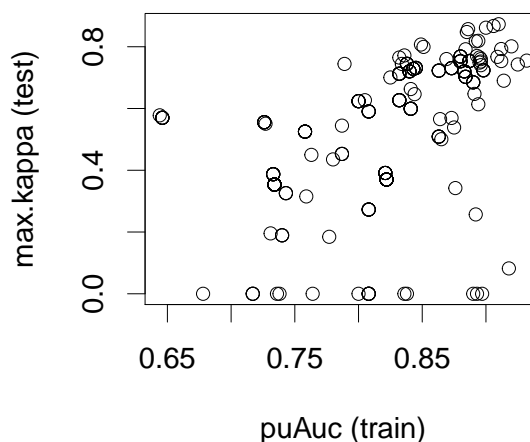

Finally, we might also want to evaluate the whole model selection process, e.g. by calculating the maximum achievable accuracy for all models. This can be useful, e.g. to compare different model selection approaches, in particular, the performance metrics and/or the resampling approaches to estimate these

176 metrics [REF]. Here for example, the best achievable κ (y-axis) of each model is plotted against the puF
177 (x-axis, left plot) and the puAuc (x-axis, right plot) metric.

```
ev.modsel <- evaluate(oc, y=te.y, u=te.x, allModels=TRUE, positive=1)
```

```
plot(ev.modsel, fromTrain='puF')
```

```
plot(ev.modsel, fromTrain='puAuc')
```



178

179

180 6 More one-class classifiers...

181 Currently three classifiers are implemented: the biased SVM [3] and one-class SVM [6] via the package
182 kernlab [7] and maxent [5] via the package dismo [8].

183 You can see the definition of the interfaces, or custom `train()` methods, with the `getModelInfoOneClass()`
184 If you want to use your own one-class classifier with the `oneClass/caret` infrastructure you can use these
185 models as examples, but see also the explanations on [http://caret.r-forge.r-project.org/custom_](http://caret.r-forge.r-project.org/custom_models.html)
186 [models.html](http://caret.r-forge.r-project.org/custom_models.html).

187 7 Parallel processing

188 Some tasks, e.g. model selection and predictions (see below), can be processed in parallel via the `foreach`
189 package. If you want to use the parallel processing capabilities you need to register a `foreach` parallel
190 backend. This can be done as follows:

```
require(foreach)
require(doParallel)
cl <- makeCluster(detectCores()-1) # leave one core free
registerDoParallel(cl)
```

8 Prediction of large raster data

If you have large raster files you might want to work with a ('rasterTiled') object. The package can also be downloaded from GitHub <https://github.com/benmack/rasterTiled>. This is a straightforward way of predicting large raster files in parallel.

The function `rasterTiled()` returns an S3 object, say `rt` which holds the raster, the valid cells and the start and end indices of the tiles.

```
require(rasterTiled)

## Loading required package: rasterTiled

rt <- rasterTiled(bananas$x)
names(rt)

## [1] "raster"      "validCells" "tiles"

rt # also estimates the approximate size of the tiles in memory

## class          : RasterBrick
## dimensions     : 400, 400, 160000, 2  (nrow, ncol, ncell, nlayers)
## resolution     : 1, 1  (x, y)
## extent        : 0, 400, 0, 400  (xmin, xmax, ymin, ymax)
## coord. ref.    : NA
## data source    : in memory
## names         :    x1,    x2
## min values     : -12.6, -11.7
## max values     : 7.941, 7.188
##
## Valid cells: 160000 of 160000.
##
## Tiles:
##          tile1 tile2  tile3  tile4
## idxStart    1 5e+04 100001 150001
## idxEnd   50000 1e+05 150000 160000
##
## Approx. size per tile (Mb):
## 10.3 10.3 10.3 2.1
```

See also the package 'spatial.tools' ...

Show an example...

9 Future work ...

- Improve PU-performance metrics
- Feature selection

- A Posteriori probabilities form the one-class output
- Efficient handling of unlabeled data: an iterative approach
- Merging one-class classification outputs of different classes

Acknowledgements

This document was produced in RStudio using the knitr package [9].

References

- [1] Max Kuhn. Contributions from Jed Wing et al. *caret: Classification and Regression Training*. R package version 6.0-24. 2014. URL: <http://CRAN.R-project.org/package=caret>.
- [2] Bing Liu et al. “Partially Supervised Classification of Text Documents”. In: 2002, pp. 387–394.
- [3] Bing Liu et al. “Building text classifiers using positive and unlabeled examples”. In: *In: Intl. Conf. on Data Mining*. 2003, pp. 179–188.
- [4] Jane Elith et al. “A statistical explanation of MaxEnt for ecologists”. In: *Diversity and Distributions* 17.1 (2011), pp. 43–57.
- [5] Steven J. Phillips and Miroslav Dudík. “Modeling of species distributions with Maxent: new extensions and a comprehensive evaluation”. In: *Ecography* 31.2 (2008), pp. 161–175.
- [6] Bernhard Schölkopf et al. “Estimating the Support of a High-Dimensional Distribution”. In: *Neural Computation* 13.7 (2001), pp. 1443–1471. ISSN: 0899-7667. DOI: 10.1162/089976601750264965.
- [7] Alexandros Karatzoglou et al. “kernlab – An S4 Package for Kernel Methods in R”. In: *Journal of Statistical Software* 11.9 (2004), pp. 1–20. URL: <http://www.jstatsoft.org/v11/i09/>.
- [8] Robert J. Hijmans et al. *dismo: Species distribution modeling*. 2013. URL: <http://CRAN.R-project.org/package=dismo>.
- [9] Yihui Xie. *knitr: A general-purpose package for dynamic report generation in R*. R package version 1.4.1. 2013. URL: <http://yihui.name/knitr/>.